

Yuvraj Soni

+91-7060642985 | yuvrajsoni2172@gmail.com | linkedin.com/in/yuvraj | github.com/yuvraj2172

Data Engineer | Big Data Engineer | Azure Data Engineer

PROFESSIONAL SUMMARY

Results-driven Data Engineer with **2+ years** of experience designing, building, and optimizing large-scale **Big Data pipelines** using **PySpark, Spark SQL, Apache NiFi, Azure Databricks, Apache Airflow**, and distributed computing frameworks on **MapR-Hadoop** and cloud platforms. Delivered **40% faster data processing** and **30% reduction in debugging time** through performance tuning, workflow automation, and observability enhancements. Expertise in **ETL/ELT pipeline development**, data lake architecture, and telecom-scale dataset processing powering **financial reporting and billing analytics**. Leverages **AI-assisted development** tools to accelerate pipeline delivery and code quality. Certified **Microsoft Fabric Data Engineer Associate (DP-700)** and **Azure Data Fundamentals (DP-900)** professional.

TECHNICAL SKILLS

Programming: Python, SQL

Big Data Processing: Apache Spark, PySpark, Spark SQL, Azure Databricks, Apache NiFi, ClickHouse, MapR-Hadoop, Distributed Computing, HDFS, Parquet

Cloud & Data Platforms: Microsoft Azure (ADLS Gen2, Data Factory, Databricks), Microsoft Fabric

Data Engineering: ETL/ELT Pipelines, Data Ingestion, Batch Processing, Data Validation, Data Quality, Data Warehouse, Partitioning Strategies

Workflow Orchestration: Apache Airflow, DAG Development, Dependency Management, Workflow Automation, GitLab CI/CD

Performance Engineering: Partition Pruning, Broadcast Joins, Caching, Shuffle Optimization, Query Optimization, Resource Tuning

Tools & Productivity: Git, Docker, Apache Superset, REST APIs, AI-Assisted Development (GitHub Copilot, LLM-based Code Review)

PROFESSIONAL EXPERIENCE

Data Engineer

Sep 2024 – Present

Ksolves India Limited

Noida, India

- Developed and deployed **scalable data pipelines** using PySpark, Spark SQL, and Python on distributed Big Data frameworks, processing **100+ GB/day** of financial and telecom datasets.
- Engineered end-to-end **ETL workflows** for data ingestion, transformation, and cleaning, powering downstream analytics dashboards in **Apache Superset**.
- Designed and optimized **Apache Spark jobs** with partition pruning, broadcast joins, and caching strategies, achieving measurable performance gains across production workloads.
- Automated **data pipeline orchestration** using Apache Airflow, reducing manual intervention, improving reliability, and increasing pipeline traceability.
- Built custom **backend data systems** for major telecom clients, enabling scalable data processing, business logic automation, and real-time reporting.
- Leveraged **AI-assisted development** tools (LLM-based code generation and review) to accelerate PySpark job development, reduce boilerplate, and improve code consistency across the team.
- Accelerated **pipeline debugging and root-cause analysis** using AI tools to interpret Spark stack traces, suggest configuration fixes, and resolve Python dependency conflicts faster.
- Maintained production-ready codebases using **Git/GitLab** for version control, branching strategies, and peer code reviews; collaborated in **Agile/Scrum** sprints including planning, backlog refinement, and cross-functional discussions.
- Interfaced directly with **US-based stakeholders** to gather requirements, translate business needs into technical specifications, and deliver high-impact analytical solutions.

Project: Enterprise Data Pipeline Modernization (Apache Drill to PySpark)

- Led full-cycle **migration of legacy data transformation workflows** from Apache Drill to PySpark, modernizing the enterprise data platform and improving overall reliability.
- Rewrote complex business logic into optimized, distributed **PySpark/Spark SQL jobs** with partition pruning, broadcast joins, and caching, cutting job runtime significantly.
- Integrated pipelines with **Apache Airflow** for scheduling, orchestration, dependency management, and real-time monitoring.

- Delivered **40% faster data processing** and **30% reduction in debugging time** via structured logging and observability improvements across dev, QA, and production.

Project: Custom Orchestration for Daily User-Level Call Processing

- Built an **end-to-end data pipeline** to ingest, process, and aggregate telecom CDRs into per-user, per-day metrics for **financial reporting and billing analysis**.
- Designed modular **PySpark jobs** (parse, validate, enrich, aggregate) with optimized transformations, partitioning, and efficient joins for high-volume call data.
- Implemented **Apache Airflow DAGs** for daily orchestration with scheduling, dependency management, retries, and monitoring – ensuring fully automated, reliable workflows.
- Standardized storage on **Parquet format** with event-date partitioning; applied **data validation** and quality checks, eliminating manual reporting for billing and revenue teams.

Project: PySpark Batch Pipeline on NiFi–ClickHouse–MapR Stack

- Sole owner of end-to-end pipeline development – designed, built, and deployed **modular PySpark batch jobs** for data ingestion from **Apache NiFi**, transformation, and loading into **ClickHouse** (columnar **data warehouse**) on a shared **MapR-Hadoop** cluster.
- Authored all **Apache Airflow DAGs** to schedule, orchestrate, and monitor PySpark job execution – configuring triggers, inter-task dependencies, retries, and SLA alerts for production reliability.
- Resolved **environment misconfigurations** and a **Python dependency conflict** on the MapR cluster independently, unblocking deployment without external escalation.
- Validated **end-to-end data integrity** across all pipeline stages pre-cutover; pipeline went live on schedule, generating **automated stakeholder reports** with zero manual intervention.

Technical Trainee – Data Engineering

Jun 2022 – Jul 2022

Gemini Solutions

Gurugram, Haryana

- Developed and tested **PySpark scripts** for processing structured and semi-structured datasets in a distributed environment.
- Performed **SQL-based data extraction**, cleaning, and validation to support downstream reporting and analytics pipelines.
- Built small-scale **batch data processing pipelines**, optimizing transformations and improving overall data quality.
- Collaborated with senior data engineers to learn production-grade **Big Data workflows**, distributed processing patterns, and data engineering best practices.

EDUCATION

GLA University

Mathura

Bachelor of Technology in Computer Science

Aug 2020 – May 2024

Gyan Deep Shiksha Bharati Sr. Sec. School

Mathura

Senior Secondary

Apr 2018 – May 2020

CERTIFICATIONS

Microsoft Certified: Fabric Data Engineer Associate (DP-700) (Microsoft, Aug 2025)

Validated proficiency in building and optimizing data pipelines, lakehouse architecture, and analytics workflows using Microsoft Fabric and Azure.

Microsoft Azure Data Fundamentals (DP-900) (Microsoft, Dec 2024)

Verified foundational knowledge of Azure cloud data concepts, storage, and analytics services including Data Factory, Synapse Analytics, and Databricks.

ACHIEVEMENTS

- Received **Spot Award (twice)** at Ksolves India Limited for delivering high-impact data engineering solutions and demonstrating exceptional technical excellence and collaboration.
- Solved **1,000+ algorithmic problems** on **LeetCode**, demonstrating strong proficiency in **Data Structures, Algorithms, and Problem-Solving**.
- Participated in multiple **Major League Hacking (MLH)** hackathons, building end-to-end **data-driven and AI-based prototypes** under tight deadlines.